

Lecture 02 Floating-point numbers and rounding

Prerequisites: Lecture 01, fractions, powers of two.

Goal: predict what a floating-point operation can preserve, lose, or certify.

1. Finite digits in a chosen base

0–12 minutes: representation by hand before inspecting Python

Writing a number with finitely many digits is a mathematical restriction. In base two, a finite fractional expansion has the form

$$0.b_1b_2\cdots b_k = \sum_{j=1}^k b_j 2^{-j} = \frac{m}{2^k}, \quad b_j \in \{0, 1\}.$$

Consequently a reduced rational a/b has a finite binary expansion exactly when b is a power of two, ignoring the machine's exponent and precision limits. If $a/b = m/2^k$ and a, b are coprime, then b divides 2^k ; conversely a power-of-two denominator directly supplies a finite expansion. In base ten the corresponding denominators may contain factors of both two and five.

Thus $3/8 = 0.011_2$ is exact in binary, while $1/10 = 0.0001100110011\dots_2$ repeats. The issue is not that a decimal literal has “too many digits”: one decimal digit can request a nonterminating binary fraction. A finite machine format must select a nearby representable value.

A toy format. With $p = 4$ significant binary digits, a positive normal number is $(1.b_1b_2b_3)_2 2^e$. In $[1, 2)$ its possibilities are $1, 1 + 1/8, 1 + 2/8, \dots, 1 + 7/8$. Moving the exponent rescales this small grid. The significand controls relative resolution; the exponent controls magnitude. Do not confuse the two roles.

Python provides two useful descriptions of a stored float:

```
x = 0.1
print(x.hex())           # 0x1.999999999999ap-4
numerator, denominator = x.as_integer_ratio()
print(numerator, denominator)
```

The hexadecimal string exposes the stored significand and a binary exponent. The ratio is the exact rational value of that bit pattern. Neither description claims that the stored value equals the intended rational $1/10$. A short printed decimal is a convenient representation of the stored float, not a transcript of every digit in its exact decimal expansion.

Warm-up: justify each answer

Which of $1/8$, $1/5$, $3/20$, and $13/16$ have finite binary expansions? Write $13/16$ in binary. Would doubling the available precision make $1/5$ exactly representable? Separate representability from the quality of approximation.

Generate the digits rather than memorizing them. For $1/10$, start with remainder $r = 1$. At each step divide $2r$ by 10: its integer quotient is the next binary digit, and its remainder becomes the new r . The remainders begin $2, 4, 8, 6, 2, \dots$; a repeated remainder restarts the same future digit sequence. This finite-state calculation explains the repetition. Apply it to $1/8$ and observe the remainder becoming zero. Terminating and repeating expansions are consequences of integer arithmetic, not quirks of Python's display.

2. Decode binary64 and derive its spacing

12–30 minutes: fields, a worked decoding, and a number line

The binary64 format has a sign bit s , an 11-bit stored exponent E , and a 52-bit fraction field interpreted as an integer F . For normal values,

$$x = (-1)^s 2^{E-1023} \left(1 + \frac{F}{2^{52}} \right), \quad 1 \leq E \leq 2046, \quad 0 \leq F < 2^{52}.$$

The leading one is implicit: there are $p = 53$ significant binary digits even though only 52 fraction bits are stored. The exponent bias is 1023.

Value	s	E	F	Normalized form
1	0	1023	0	$1.0_2 2^0$
1.5	0	1023	2^{51}	$1.1_2 2^0$
-2.5	1	1024	2^{50}	$-1.01_2 2^1$

For -2.5 , divide the magnitude by two to obtain $1.25 = 1.01_2$. The unbiased exponent is one, and the fractional significand is $1/4$; therefore $F/2^{52} = 1/4$. This derives the fields rather than guessing their bits.

Spacing theorem. In a normal binade $[2^e, 2^{e+1})$, consecutive values are separated by $2^{e-(p-1)}$. Indeed, increasing F by one changes the significand by $2^{-(p-1)}$; multiplication by 2^e gives the stated step. The spacing doubles at the next binade. Below an exact power of two, the preceding normal binade has half the spacing of the one above.

A toy format with four significant binary digits



```
import math
for x in [1.0, 2.0, 2.0**53]:
    below = math.nextafter(x, -math.inf)
    above = math.nextafter(x, math.inf)
    print(x, x - below, above - x)
```

At 1, the lower and upper gaps are 2^{-53} and 2^{-52} . At 2^{53} , they are 1 and 2. All integers from 0 through 2^{53} are representable, but not all larger integers are. Precision is not a fixed number of decimal places: absolute resolution changes with scale.

Predict before running

Why can $2^{53} - 1$ be represented when $2^{53} + 1$ cannot? What are the two gaps at 2^{54} ? The smallest normal number is a special boundary; page 4 explains why.

The largest finite normal value is $(2 - 2^{-52})2^{1023}$: use the largest normal exponent and set every fraction bit to one. This limit concerns *range*, whereas spacing concerns *resolution*. Increasing the exponent range alone would not let us distinguish 2^{53} from $2^{53} + 1$ at 53-bit precision.

3. Rounding modes and a one-operation bound

30–48 minutes: prove the bound, then test midpoint cases

Rounding to nearest selects a representable value closest to the exact result. At an exact midpoint, *ties to even* selects the neighbour whose last significand bit is zero. It does not mean rounding every halfway value toward zero. Directed rounding instead selects the neighbouring representable value below or above the exact result. Those two modes will support interval endpoints later.

We use two distinct constants:

$$\varepsilon = 2^{1-p} = 2^{-52} \quad (\text{gap above } 1), \quad u = 2^{-p} = 2^{-53} \quad (\text{unit roundoff}).$$

Some sources call either quantity “machine epsilon”, so a symbol without a definition is not enough. Our convention fixes $\varepsilon = 2u$.

Relative rounding bound. For a nonzero exact value z in the normal range, assuming round-to-nearest and no overflow,

$$\text{fl}(z) = z(1 + \delta), \quad |\delta| \leq u.$$

To prove this for $z > 0$, take $2^e \leq z < 2^{e+1}$. The maximum distance to a nearest neighbour is half a spacing, 2^{e-p} . Dividing by $z \geq 2^e$ gives $|\text{fl}(z) - z|/z \leq 2^{-p}$. Negative values follow by symmetry. The normal-range assumption matters; it will fail in our next example.

For a correctly rounded basic operation, apply this statement to the exact result of that operation on its *stored operands*. Errors already present in the operands have not disappeared. Nor does a sequence of a thousand operations acquire a global relative-error bound of u merely because each local rounding has such a bound.

```
u = 2.0**-53
print(1.0 + u == 1.0)           # True: midpoint, even lower neighbour
print(1.0 + 3*u == 1.0 + 4*u)  # True: even upper neighbour
print(2.0**53 + 1.0 == 2.0**53) # True: an exact increment is lost
```

In the last line both input operands are exact. The exact sum lies halfway between 2^{53} and $2^{53} + 2$; the first has the even low significand bit. This is arithmetic rounding, not an inaccurate representation of the integer one.

Small-group proof check

Explain each inequality in the rounding-bound proof. Where were normal range, nearest rounding, and exact operands used? If the exact result is zero, how would you state an absolute-error claim without dividing by zero?

Preview. To enclose a compound expression, it is not enough to round its final printed value outwards. Intermediate errors and dependencies also need control. Interval arithmetic will organize those operations and their assumptions.

A first accumulated-error calculation. For stored inputs a, b, c , assume the rounding model applies to both multiplications. Then

$$\text{fl}(\text{fl}(ab)c) = abc(1 + \delta_1)(1 + \delta_2), \quad |\delta_1|, |\delta_2| \leq u.$$

The total relative error, when $abc \neq 0$, is at most $2u + u^2$, not u . This follows by expanding the product and applying the triangle inequality. Even this tiny example needs a composed argument; cancellation in more general expressions will make the analysis more delicate.

4. The boundary of the model

48–62 minutes: subnormals, exceptional values, and input provenance

For normal binary64 numbers, the unbiased exponent ranges from -1022 to 1023 . The smallest positive normal value is 2^{-1022} . Exponent field $E = 0$ uses *subnormal* values instead:

$$x = (-1)^s F 2^{-1074}, \quad 1 \leq F < 2^{52}.$$

These have no implicit leading one. Their fixed spacing $\alpha = 2^{-1074}$ fills the gap towards zero. The gap just below the smallest normal equals the gap just above it, unlike a boundary between two normal binades. As magnitudes decrease through the subnormal region, relative precision deteriorates.

For example, the exact real $z = \alpha/2$ rounds to zero under ties to even. Its absolute error is only $\alpha/2$, but its relative error is one. A useful combined bound, excluding overflow, is

$$|\text{fl}(z) - z| \leq u|z| + \alpha/2.$$

Normal results are covered by the relative estimate; in the tiny-value region, nearest rounding has absolute error at most half the fixed spacing. The extra absolute term accounts for gradual underflow.

```
from fractions import Fraction
small = Fraction(1, 2**1075)
rounded = float(small)
assert rounded == 0.0
assert abs(Fraction.from_float(rounded) - small) / small == 1
```

Special encodings. $E = 0, F = 0$ represents signed zero. With $E = 2047, F = 0$ represents an infinity and $F \neq 0$ a NaN. Infinities and NaNs are not ordinary real-number answers. In particular, `nan == nan` is false. Python scalar division by zero raises an exception; NumPy arithmetic can instead produce infinities or NaNs with warnings. A proof must specify the actual operation semantics rather than assuming every programming interface behaves identically.

More precision cannot repair the meaning of an input.

```
from decimal import Decimal, localcontext
with localcontext() as context:
    context.prec = 60
    intended = Decimal("0.1")
    inherited = Decimal(0.1)
    print(inherited - intended)
    print(Decimal(1) / Decimal(3))
```

The string denotes the exact decimal tenth; the float constructor preserves the already-rounded binary input. The third still needs rounding at finite decimal precision. Increasing the context changes later arithmetic, not earlier input history.

5. A time-conversion error that grows with uptime

62–77 minutes: exact model, operational bound, and historical context

The GAO report on the 1991 Patriot failure describes precision loss when integer clock ticks, counted in tenths of seconds, were converted to time. The tracking error grew with operating duration; Appendix II reports about 0.3433 seconds at 100 hours. This was not a loop repeatedly adding Python’s binary64 0.1. The incident motivates a question: over what operating range is an error bound acceptable? [GAO IMTEC-92-26, pp. 5–6 and Appendix II]

Use the following *simplified fixed-point model*, not a reconstruction of the historical software. Approximate $c = 1/10$ seconds per tick by truncation to q fractional binary bits:

$$\hat{c} = \frac{\lfloor 2^q/10 \rfloor}{2^q}, \quad 0 \leq c - \hat{c} < 2^{-q}.$$

For $N \geq 0$ exact integer ticks, compare $t = Nc$ with $\hat{t} = N\hat{c}$. In this model their difference is exact:

$$\Delta t = N(c - \hat{c}), \quad 0 \leq \Delta t < N2^{-q}.$$

There is no random cancellation of the per-tick bias. Additional rounding of a real implementation’s multiplication would need another term; the model above isolates the conversion-coefficient error.

```
from fractions import Fraction
q = 23
scale = 2**q
tick = Fraction(1, 10)
stored_tick = Fraction(scale // 10, scale)
error_per_tick = tick - stored_tick
for hours in [1, 8, 24, 100]:
    ticks = hours * 3600 * 10
    error = ticks * error_per_tick
    print(hours, error, float(error))
```

Here $c - \hat{c} = 1/10485760$ seconds. At 100 hours, $N = 3600000$ and $\Delta t = 5625/16384 \approx 0.343322754$ seconds. Matching the scale of the historical table does not establish equivalence with all details of the historical implementation.

Design from a tolerance. To guarantee a model error below τ over $N_{\max}$ ticks, it suffices to choose q with $N_{\max}2^{-q} \leq \tau$. Round-to-nearest gives a coefficient-error bound of 2^{-q-1} instead. If the true tick length is uncertain by at most η , add $N\eta$ to the conversion bound. More coefficient bits do not remove physical clock uncertainty.

Design exercise

For 100 hours and $\tau = 0.01$ seconds, find a sufficient q using the general truncation bound. Explain why an uptime limit belongs in the specification.

6. Exercises and a rounding audit

77–90 minutes: compare explanations; complete remaining calculations after class

Use exact fractions for reference values whenever possible. State whether you are discussing an input conversion, one arithmetic operation, or an entire algorithm.

1. **Decode a value.** Write 6.5 in normalized binary form and give its binary64 fields s, E, F . Find its neighbours' spacing. Repeat the spacing question at 8 and explain why the lower and upper gaps now differ.
2. **Tie-breaking.** Work in a toy binary format with $p = 4$ and sufficient exponent range. Round $1 + 1/16$, $1 + 3/16$, and $2 + 1/8$ to nearest, ties to even. For each, name both neighbours and identify the even significand. Do not use a binary64 calculation as a substitute for this toy-format calculation.
3. **A lost integer.** Predict the binary64 results of adding 1, 2, and 3 to 2^{54} . Explain why subtracting 2^{54} afterwards cannot recover the discarded information. Which additions involve a midpoint tie?
4. **An input audit.** Compare `Fraction("0.1")` with `Fraction(0.1)`. Then compare `Decimal("0.1")` with `Decimal(0.1)`. Which preserve the intended decimal tenth, and which preserve the stored float? Would setting a decimal precision of 100 change these input meanings?
5. **Bound the time error.** Derive the page-5 bound directly from the floor inequality. For $q = 23$, give a sufficient maximum integer tick count that keeps the general bound at most 0.1 seconds. Compare it with what is permitted using the exact coefficient error. Explain why agreement at one uptime is not a proof about all possible runtimes or implementations.

Selected checks. Exercise 1: $6.5 = 1.101_2 2^2$, so $s = 0$, $E = 1025$, and $F = 5 \cdot 2^{49}$; its spacing is 2^{-50} . At 8, the lower gap is 2^{-50} and the upper gap 2^{-49} . Exercise 2 gives 1, 1.25, and 2. Exercise 3 gives 2^{54} , 2^{54} , and $2^{54} + 4$. For the page-5 design exercise, $q = 29$ is sufficient. Check each answer by explaining its rounding or bound, not just by matching the number.

Exit ticket

State the meanings of ε and u . Give one hypothesis needed for $\text{fl}(z) = z(1 + \delta)$ with $|\delta| \leq u$, and a concrete counterexample when that hypothesis is removed. Which mistake would persist even at 200-digit precision?

Reading and notebook. Companion: `notebooks/02_floating_point.ipynb`; Lab 02: inspect the machine. Tucker (2011), §§1.2–1.5. See the [Python floating-point tutorial](#) and [Decimal documentation](#) for Python interfaces. Historical source: GAO IMTEC-92-26 (1992), linked on page 5.